

Live Model Inference

Run large models directly from the dashboard with real-time metrics.

GPU Live Telemetry

Refresh GPU

☐ Auto-refresh

Refresh interval (sec)

2

GPU Util

0.0%

VRAM

2597/40...

↑ 63.4%

Temp

69.0 C

Power

15.3 W

↑ limit N/A

CUDA Allocated

8 MB

Max CUDA Allocation

4096 MB


↑ device total VRAM

Energy (window)

2.9355 Wh

Source: nvml | GPU: NVIDIA GeForce RTX 3050 Ti Laptop GPU | Samples: 600

Live GPU Trends





Export Telemetry History CSV

☒ GPU-Only Inference Mode ?

CUDA is available. Inference will use GPU.

1 Select or Configure Model

Popular Models

mistralai/Mistral-7B-v0.1

Preload Model (Warmup)

2 Generation Parameters

Prompt

Model: mistralai/Mistral-7B-v0.1

Mode: Local (GPU if available)

Explain Express.Js in simple terms.

Max Tokens

91

Temperature

0.70

Top-p

Top-k

0.90

50

-

+

Advanced Options

Enable Offloading

Use KV Cache

Model Dtype

float16

Device

cuda:0

Cache Precision

INT8

Offload Dir

/tmp/kai_swap

Run Mode

baseline

3 Generation

Output appears in: 'Generated Output' and also in 'Prompt Run History' below.

Generate

Stop

Clear Model Cache

Generated Output

Latest Generation

Explain Express.Js in simple terms.

Express.js is a web application framework for Node.js, released as free and open-source software under the MIT License. It is designed for building web applications and APIs. It is used to create server-side web applications, which handle HTTP requests and responses.

Generation Metrics

It provides a set of features and tools that make it easier to build and maintain web applications. It includes features such as routing, middleware, templating

Duration	Tokens	Speed	Device	KV Cache
4340.51s	91	0.0 tok/s	cuda:0	X Disabled

Model cache hit: no | KV runtime: baseline_generate

History source: CSV + session history

Current model runtime cache: miss | Prompt stored in CSV: yes

Energy (Wh)	Tokens/Wh	Efficiency Level
2.4599	36.99	Good

Baseline mode: runtime cache disabled.

Baseline mode active: KAI KV reuse disabled.

Prompt Run History

Total Runs

8

Clear Run History

Loaded 8 runs from CSV + session history

Select a run

6. 2026-05-16 20:57:48 | kai | completed | mistralai/Mistral-7B-v0.1 tok



Mode: KAI | Model: Express.js is a web application framework for Node.js. It is used to build web applications and APIs. It is a minimal and flexible Node.js web application framework, providing a robust set of features for web and mobile applications.

Express.js is built on top of Node.js and provides a set of features that make it easy to build web applications. It includes a powerful HTTP server, a routing system, | Device: 1904.2106211185455 | KV: OFF (False)

KV runtime mode: INT8

KV low-level: hit=baseline_generate miss=0 | reused_prefix_tokens=0 | new_prefill_tokens=0

Prompt

baseline

Response

Explain Express.Js in simple terms.

Express.js is a web application framework for Node.js. It is used to build web applications and APIs. It is a minimal and flexible Node.js web application framework, providing a robust set of features for web and mobile applications.

Express.js is built on top of Node.js and provides a set of features that make it easy to build web applications. It includes a powerful HTTP server, a routing system,

KAI vs Baseline Compare

Latest KAI Run

Duration

2017.41s

Tokens/sec

0.05

Energy (Wh)

Latest Baseline Run

Duration

4340.51s

Tokens/sec

0.02

Energy (Wh)

2.8184
Tokens/Wh
32.29

2.4599
Tokens/Wh
36.99

Latest Baseline vs KAI

