

Live Model Inference

Run large models directly from the dashboard with real-time metrics.

GPU Live Telemetry

Refresh GPU

☐ Auto-refresh

Refresh interval (sec)

2

GPU Util

VRAM

Temp

Power

CUDA Allocated

Max CUDA Allocation

Energy (window)

0.0%

2909/...

78.0 C

31.3 W

8 MB

4096 MB

1.8380...


↑ 71.0%

↑ limit N/A

↑ device total VRAM

Source: nvml | GPU: NVIDIA GeForce RTX 3050 Ti Laptop GPU | Samples: 144

Live GPU Trends



Export Telemetry History CSV

☒ GPU-Only Inference Mode ⓘ

CUDA is available. Inference will use GPU.

1 Select or Configure Model

Popular Models

microsoft/phi-2


Preload Model (Warmup)

2 Generation Parameters

Prompt

Explain quantum computing in simple terms.

Model: microsoft/phi-2
Mode: Local (GPU if available)



Advanced Options

☒ Enable Offloading

☒ Use KV Cache

Model Dtype

float16

Device

cuda:0

Cache Precision

INT8

Offload Dir

/tmp/kai_swap

Run Mode

baseline

3 Generation

Output appears in: 'Generated Output' and also in 'Prompt Run History' below.

Generate

Stop

Clear Model Cache

Generated Output

Latest Generation

Explain quantum computing in simple terms. Quantum computing is a type of computing that uses the principles of quantum mechanics to perform calculations. Unlike classical computers, which use bits that can be either 0 or 1, quantum computers use quantum bits, or qubits, that can be both 0 and 1 at the same time. This allows quantum computers to perform multiple calculations simultaneously, making them much faster than classical computers. Quantum computing is still in its early stages, but it has the potential to revolutionize fields such as cryptography, drug discovery, and optimization problems.

Generation Metrics

Duration	Tokens	Speed	Device	KV Cache
85.92s	104	1.2 tok/s	cuda:0	X Disabled

Model cache hit: no | KV runtime: baseline_generate

History source: CSV + session history

Current model runtime cache: miss | Prompt stored in CSV: yes

Energy (Wh)

0.8077

Tokens/Wh

128.75

Efficiency Level

Excellent

Baseline mode: runtime cache disabled.

Baseline mode active: KAI KV reuse disabled.

Prompt Run History

Total Runs

4

Clear Run History

Loaded 4 runs from CSV + session history

Select a run

3. 2026-05-16 19:48:27 | kai | completed | 100 tok

Mode: BASELINE | Model: microsoft/phi-2 | Device: cuda:0 | KV: OFF (INT8)

KV runtime mode: baseline_generate

KV low-level: hit=0 miss=0 | reused_prefix_tokens=0 | new_prefill_tokens=8

Prompt

Explain quantum computing in simple terms.

Response

Quantum computing is a type of computing that uses the principles of quantum mechanics to perform calculations. Unlike classical computers, which use bits that can be either 0 or 1, quantum computers use quantum bits, or qubits, that can be both 0 and 1 at the same time. This allows quantum computers to perform multiple calculations simultaneously, making them much faster than classical computers. Quantum computing is still in its early stages, but it has the potential to revolutionize fields such as cryptography, drug discovery, and optimization problems.

KAI vs Baseline Compare

Latest KAI Run

Duration

83.11s

Tokens/sec

1.20

Energy (Wh)

0.7364

Latest Baseline Run

Duration

85.92s

Tokens/sec

1.21

Energy (Wh)

0.8077

Tokens/Wh

135.79

Tokens/Wh

128.75

Latest Baseline vs KAI

